

A thick dark blue vertical bar is positioned on the left side of the page. To its right, several thin, curved lines in shades of blue and grey sweep upwards and outwards from the bottom left corner.

Knowledge test

ARCHANA G S
NSTI CALICUT

ACTIVITY 1

AIM :- Building a Simple Neural Network (10 marks) Build and compile a simple neural network using Keras to classify the MNIST dataset (handwritten digits). The model should include at least one hidden layer. Provide the code and briefly explain each step.

Requirements :-

- Computer
- Vs code
- Network

Procedure :-

1. Create a folder
2. Open vs code
3. Create a py file in that folder
4. Write the code in that file
5. Import Libraries:

TensorFlow and Keras libraries are imported for building the model.

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.datasets import mnist
```

6. Load Dataset:

The MNIST dataset is loaded and normalized to the range [0, 1].

```
# Load the MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0 # Norm
```

7. Build Model:

Flatten Layer: Converts the 28x28 images into a 1D array of 784 pixels.

Dense Layer: A hidden layer with 128 neurons and ReLU activation function.

Output Layer: A dense layer with 10 neurons (one for each digit) and softmax activation for multi-class classification.

```
# Build the model
model = models.Sequential()
model.add(layers.Flatten(input_shape=(28, 28)))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

8. Compile Model:

The model is compiled using the Adam optimizer and sparse categorical cross-entropy loss.

```
# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

9. Train Model:

The model is trained on the training data for 5 epochs.

```
# Train the model
model.fit(x_train, y_train, epochs=5)
```

10. Evaluate Model:

The model's performance is evaluated on the test dataset, and the accuracy is printed.

```
# Evaluate the model
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_acc}')
```

Output

```
2024-08-02 17:05:58.612229: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is
ns in performance-critical operations.
To enable the following instructions: AVX2 AVX512F AVX512_VNNI FMA, in other operations, rebuild TensorFlow
Epoch 1/5
1875/1875 ██████████ 4s 2ms/step - accuracy: 0.8792 - loss: 0.4360
Epoch 2/5
1875/1875 ██████████ 5s 2ms/step - accuracy: 0.9644 - loss: 0.1197
Epoch 3/5
1875/1875 ██████████ 3s 2ms/step - accuracy: 0.9762 - loss: 0.0780
Epoch 4/5
1875/1875 ██████████ 5s 2ms/step - accuracy: 0.9829 - loss: 0.0566
Epoch 5/5
1875/1875 ██████████ 4s 2ms/step - accuracy: 0.9859 - loss: 0.0446
313/313 ██████████ 0s 1ms/step - accuracy: 0.9721 - loss: 0.0890
Test accuracy: 0.9761999845504761
PS C:\Users\User\Desktop\OJT\02.08.2024> █
```

ACTIVITY 2

AIM :- Data Augmentation (5 marks) Implement data augmentation on a given image dataset using Keras. Show at least three different augmentation techniques and explain how they help improve model performance

Requirements :-

- Computer
- Vs code
- Network

Procedure :-

1. Create a folder
2. Open vs code
3. Create a py file in that folder
4. Copy a image for data augmentaion
5. Write the code in that file
6. Import Libraries:
Necessary libraries for image processing and visualization are imported.

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
import os
```

7. Check Image Path: The code checks if the specified image file exists.

```
# Define the path to your
image_path = 'images.jpg'
```

8. ImageDataGenerator:

An instance is created with various augmentation techniques

```
# Check if the file exists
if not os.path.isfile(image_path):
    raise FileNotFoundError(f"Image file {image_path} does not exist")

# Create an instance of ImageDataGenerator
datagen = ImageDataGenerator(
    rotation_range=40,          # Random rotation
    width_shift_range=0.2,      # Random width shift
    height_shift_range=0.2,     # Random height shift
    shear_range=0.2,           # Random shear
    zoom_range=0.2,            # Random zoom
    horizontal_flip=True,       # Random horizontal flip
    fill_mode='nearest'        # Fill missing pixels
)
```

9. Load and Preprocess Image:

The image is loaded and converted into a batch format.

```
# Load and preprocess the image
image = tf.keras.preprocessing.image.load_img(image_path)
image = tf.keras.preprocessing.image.img_to_array(image)
image = np.expand_dims(image, axis=0) # Convert image to a batch
```

10. Apply Augmentations: Augmented images are generated using the flow method.

```
# Apply augmentations
augmented_images = datagen.flow(image, batch_size=1)
```

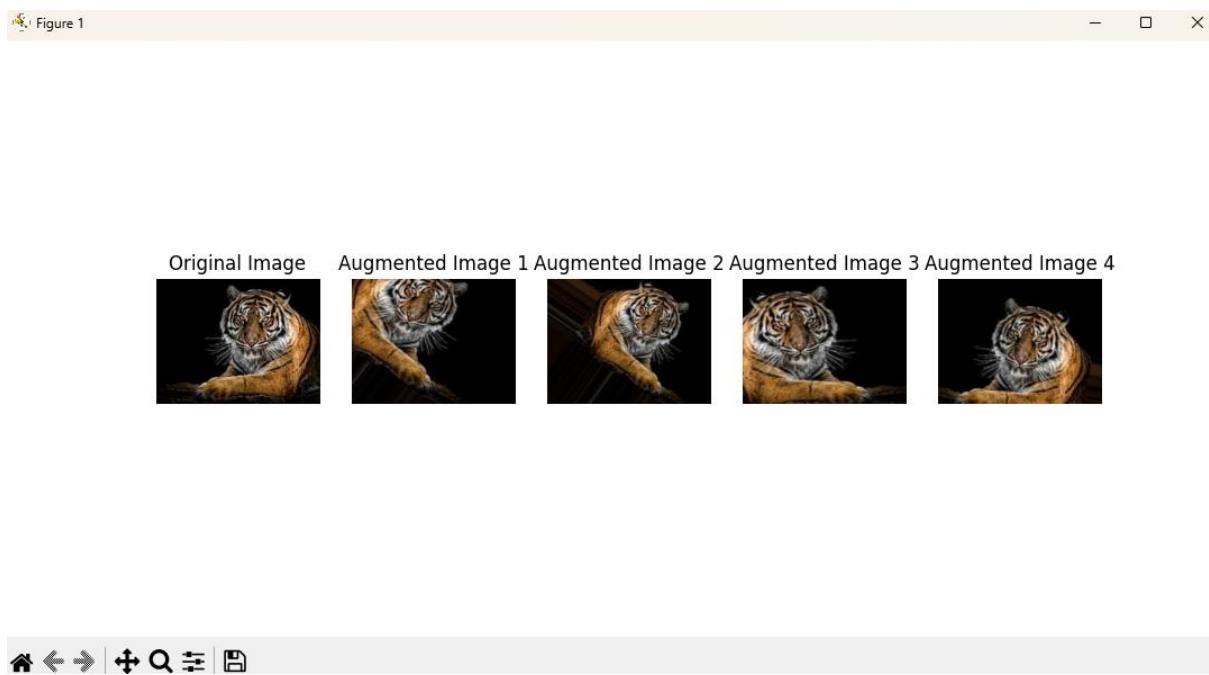
10. Plot Images:

The original and augmented images are displayed for comparison.

```
# Plot the original and augmented images
plt.figure(figsize=(15, 15))

# Plot the original image
plt.subplot(1, 5, 1)
plt.imshow(image[0].astype('uint8'))
plt.title('Original Image')
plt.axis('off')
```

OUTPUT



ACTIVITY 3

AIM :- Custom Loss Function (5 marks) Implement a custom loss function in TensorFlow/Keras. Explain the purpose of the loss function and provide an example scenario where it would be useful.

Requirements :-

- Computer
- Vs code
- Network

Procedure :-

- 1.Create a folder
- 2.Open vs code
- 3.Create a py file in that folder
- 4.Write the code in that file
5. Import Libraries:

TensorFlow and Keras libraries are imported.

```
import tensorflow as tf
from tensorflow.keras.losses import Loss
```

2. Define Custom Loss Class: A class CustomLoss is defined, inheriting from Keras' Loss class.


```
# Define a custom loss function
Codeium: Refactor | Explain
class CustomLoss(Loss):
    Codeium: Refactor | Explain | Generate Docstring | X
    def __init__(self, threshold=1.0, **kwargs):
        super().__init__(**kwargs)
        self.threshold = threshold

    Codeium: Refactor | Explain | Generate Docstring | X
    def call(self, y_true, y_pred):
        # Compute the absolute error
        absolute_error = tf.abs(y_true - y_pred)
        # Apply custom penalty: increase loss if error exceeds threshold
        custom_loss = tf.where(absolute_error > self.threshold,
                               2 * absolute_error,
                               absolute_error)
        return tf.reduce_mean(custom_loss)
```

3. Model Definition:

A simple sequential model is created with a dense layer.

```
# Example Usage
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(10, activation='relu', input_shape=(5,)),
    tf.keras.layers.Dense(1)
])

model.compile(optimizer='adam', loss=CustomLoss(alpha=0.5))
```

4. Train Model: The model is trained for 5 epochs.

```
# Train the model
history = model.fit(X_train, y_train, epochs=5)
```

Output

2024-08-02 17:24:10.525392: I tensorflow/core/util/port.cc:113] oneDNN custom operations are on. You may see slightly different results from floating-point round-off errors from different computation orders. To turn them off, set the environment variable TF_DISABLE_ONE_DNN_ROUNDING to 1. To enable the following instructions: AVX2 AVX512F AVX512_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

Epoch 1/5
4/4 ————— 1s 4ms/step - loss: 0.4749
Epoch 2/5
4/4 ————— 0s 0s/step - loss: 0.4751
Epoch 3/5
4/4 ————— 0s 3ms/step - loss: 0.4325
Epoch 4/5
4/4 ————— 0s 2ms/step - loss: 0.4205
Epoch 5/5
4/4 ————— 0s 0s/step - loss: 0.4101
Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 10)	110
dense_1 (Dense)	(None, 1)	11

Total params: 365 (1.43 KB)
Trainable params: 121 (484.00 B)
Non-trainable params: 6 (0.00 B)
Optimizer params: 244 (980.00 B)
Training history:
'loss': [0.493704617023468, 0.4624812602996826, 0.43306484818458557, 0.4153471291065216, 0.39721455518458557]
PS C:\Users\User\Desktop\OJT\02.08.2024>

Battery saver

Battery saver is on

ACTIVITY 4

AIM :- Transfer Learning (5 marks) Use a pre-trained model (such as VGG16 or ResNet) available in Keras for a simple image classification task. Fine-tune the model for a new dataset and describe the steps taken

Requirements :-

- Computer
- Vs code
- Network

Procedure

- 1.Create a folder name as exam
- 2.Open vs code
- 3.Create a py file in that folder
- 4.Write the code in that file
- 5.Import Necessary Libraries

```
import tensorflow as tf
from tensorflow.keras.applications import VGG16
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
from tensorflow.keras.optimizers import Adam
import matplotlib.pyplot as plt
```

- 6.give paths

```
# Paths to your dataset directories
train_dir = 'C:/Users/USER/Desktop/exam/flower'
validation_dir = 'C:/Users/USER/Desktop/exam/flower'
```

- 7.create an imagedatagenerator for data augmentation

```
# Create an ImageDataGenerator for data augmentation
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

test_datagen = ImageDataGenerator(rescale=1./255)
```

8.load data

```
# Load data
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(224, 224), # Adjust based on model input size
    batch_size=32,
    class_mode='categorical'
)

validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(224, 224), # Adjust based on model input size
    batch_size=32,
    class_mode='categorical'
)
```

10. Load the Pre-Trained Model: Load the pre-trained VGG16 model without the top layers.

```
# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Normalize the images to the range [0, 1]
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Convert class vectors to binary class matrices
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)
```

✓ 1m 11.4s

11. Freeze the Pre-Trained Layers: Prevent the pre-trained layers from being updated during training.

```
# Freeze the layers of the base model
for layer in base_model.layers:
    layer.trainable = False
```

12. Add Custom Layers: Add new layers to adapt the model to the new dataset

```
# Add custom layers
x = base_model.output
x = Flatten()(x)
x = Dense(512, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# Create the final model
model = Model(inputs=base_model.input, outputs=predictions)
```

13. Compile the Model: Define the optimizer, loss function, and evaluation metrics. Train the model on the new dataset. Assess the performance of the model on the test set.

```
# Create the final model
model = Model(inputs=base_model.input, outputs=predictions)

model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])

model.fit(x_train, y_train, epochs=10, batch_size=32, validation_data=(x_test, y_test))
```

Output

```
1563/1563 — 60s 38ms/step - accuracy: 0.6213 - loss: 1.0790 - val_accuracy: 0.5890 - val_loss: 1.1612
Epoch 4/10
1563/1563 — 59s 38ms/step - accuracy: 0.6411 - loss: 1.0208 - val_accuracy: 0.6081 - val_loss: 1.1174
Epoch 5/10
1563/1563 — 64s 41ms/step - accuracy: 0.6605 - loss: 0.9668 - val_accuracy: 0.6172 - val_loss: 1.0930
Epoch 6/10
1563/1563 — 58s 37ms/step - accuracy: 0.6718 - loss: 0.9255 - val_accuracy: 0.6195 - val_loss: 1.1011
Epoch 7/10
1563/1563 — 62s 40ms/step - accuracy: 0.6946 - loss: 0.8747 - val_accuracy: 0.6088 - val_loss: 1.1394
Epoch 8/10
1563/1563 — 60s 38ms/step - accuracy: 0.7045 - loss: 0.8397 - val_accuracy: 0.6104 - val_loss: 1.1395
Epoch 9/10
1563/1563 — 64s 41ms/step - accuracy: 0.7225 - loss: 0.7831 - val_accuracy: 0.6172 - val_loss: 1.1522
Epoch 10/10
1563/1563 — 59s 38ms/step - accuracy: 0.7366 - loss: 0.7540 - val_accuracy: 0.6186 - val_loss: 1.1569
```

```
... 313/313 — 10s 31ms/step - accuracy: 0.6166 - loss: 1.1521
Test accuracy: 0.6186000108718872
```